

# Tutorial 05: Two Sample Confidence Intervals

```
#| edit: false
#| output: false
webr::install("gradethis", quiet = TRUE)
library(gradethis)
options(webr.exercise.checker = function(
  label, user_code, solution_code, check_code, envir_result, evaluate_result,
  envir_prep, last_value, engine, stage, ...
) {
  if (is.null(check_code)) {
    # No grading code, so just skip grading
    invisible(NULL)
  } else if (is.null(label)) {
    list(
      correct = FALSE,
      type = "warning",
      message = "All exercises must have a label."
    )
  } else if (is.null(solution_code)) {
    list(
      correct = FALSE,
      type = "warning",
      message = htmltools::tags$div(
        htmltools::tags$p("A problem occurred grading this exercise."),
        htmltools::tags$p(
          "No solution code was found. Note that grading exercises using the ",
          htmltools::tags$code("gradethis"),
          "package requires a model solution to be included in the document."
        )
      )
  } else {
    gradethis::gradethis_exercise_checker(
```

```

        label = label, solution_code = solution_code, user_code = user_code,
        check_code = check_code, envir_result = envir_result,
        evaluate_result = evaluate_result, envir_prep = envir_prep,
        last_value = last_value, stage = stage, engine = engine)
    }
})

```

## Q1 — Pooled (Equal Variances): Reviews

Using the local file recipe\_reviews.csv, compute a 95% two-sided confidence interval for

$$\mu_{FiveStar} - \mu_{NotFive}$$

where the outcome is best\_score and the groups are stars == 5 (FiveStar) vs stars != 5 (NotFive). Read the CSV with `read.csv("recipe_reviews.csv")`, build the two groups, and print the 2-element vector `c(lower, upper)`.

Photo by Maximus Mazar on Unsplash

### Info

For a pooled two-sample CI with equal variances:

$$\bar{x}_1 - \bar{x}_2 \pm t_{df=n_1+n_2-2}^* s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}, \quad s_p^2 = \frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2}$$

Use  $\alpha = 0.05 \implies t^* = t_{1-\frac{\alpha}{2}, df}$

Group 1 = FiveStar (rows with stars == 5), Group 2 = NotFive (rows with stars != 5).

### Preview

The “Recipe Reviews and User Feedback Dataset” is a comprehensive repository of data encompassing various aspects of recipe reviews and user interactions. It includes essential information such as the recipe name, its ranking on the top 100 recipes list, a unique recipe code, and user details like user ID, user name, and an internal user reputation score. Each review comment is uniquely identified with a comment ID and comes with additional attributes, including the creation timestamp, reply count, and the number of up-votes and down-votes received. Users’ sentiment towards recipes is quantified on a 1 to 5 star rating scale, with a score of 0 denoting an absence of rating. This dataset is a valuable resource for researchers and data scientists, facilitating endeavors in sentiment analysis, user behavior analysis, recipe recommendation systems, and more. It offers a window into the dynamics of recipe reviews and user feedback within the culinary website domain.

Run this code chunk to get a glimpse of the dataset. Feel free to change the values to visualize more/less number of rows.

```
#| echo: true
df <- read.csv("recipe_reviews.csv")
head(df[, c("recipe_name", "stars", "best_score", "thumbs_up", "thumbs_down", "user_reputation")]
df2 <- subset(df, is.finite(best_score) & is.finite(stars))
df2$group <- ifelse(df2$stars == 5, "FiveStar", "NotFive")
set.seed(258)
grp_idx <- split(seq_len(nrow(df2)), df2$group)
idx <- unlist(lapply(grp_idx, function(i) sample(i, min(200, length(i)))))
df3 <- df2[idx, ]
ylim <- quantile(df3$best_score, c(.01,.99), na.rm = TRUE)
par(mar = c(4,4,1,1))
boxplot(best_score ~ group, data = df3, notch = FALSE, outline = FALSE, ylim = ylim,
ylab = "best_score", xlab = "Group")
```

```
#| exercise: q1_pooled_reviews
#| exercise.lines: 10
#| echo: false
df <- read.csv("recipe_reviews.csv")

x1 <- subset(df, stars == 5)$ #segregate the dataset based on best_scores
x2 <- subset(df, stars != 5)$

n1 <- sum(!is.na(x1)); n2 <- sum(!is.na(x2))
m1 <- (x1, na.rm = TRUE) #compute mean of x1,x2
m2 <- (x2, na.rm = TRUE)
s1 <- (x1, na.rm = TRUE)
s2 <- (x2, na.rm = TRUE) #compute standard deviation of x1,x2

dfree <- #insert degrees of freedom
sp2 <- #calculate $s^2_p$
sp <- sqrt(sp2)

tstar <- #find critical value
se <- #compute the standard error using formula given above
```

Confirm you're comparing the means of two groups formed by a boolean condition on stars. Use the pooled-variance standard error and the appropriate two-sided t critical value for a 95% interval. Print just the lower and upper bounds as a numeric vector.

*Solution.*

```
#| exercise: q1_pooled_reviews
#| solution: true
df <- read.csv("recipe_reviews.csv")

x1 <- subset(df, stars == 5)$best_score
x2 <- subset(df, stars != 5)$best_score

n1 <- sum(!is.na(x1)); n2 <- sum(!is.na(x2))
m1 <- mean(x1, na.rm = TRUE); m2 <- mean(x2, na.rm = TRUE)
s1 <- stats::sd(x1, na.rm = TRUE); s2 <- stats::sd(x2, na.rm = TRUE)

dfree <- n1 + n2 - 2
sp2 <- (((n1 - 1) * s1^2) + ((n2 - 1) * s2^2)) / dfree
sp <- sqrt(sp2)

tstar <- qt(0.975, dfree)
se <- sp * sqrt(1/n1 + 1/n2)
est <- m1 - m2

c(est - tstar*se, est + tstar*se)

#| exercise: q1_pooled_reviews
#| check: true
gradethis::grade_this({
  df <- tryCatch(read.csv("recipe_reviews.csv"), error = function(e) NULL)
  if (is.null(df)) fail("Couldn't read 'recipe_reviews.csv'.")

  need <- c("best_score", "stars")
  if (!all(need %in% names(df))) fail("CSV must contain 'best_score' and 'stars' columns.")

  # Build groups (work even if stars is character or numeric)

  x1 <- df$best_score[df$stars == 5]
  x2 <- df$best_score[df$stars != 5]

  if (sum(is.finite(x1)) < 2 || sum(is.finite(x2)) < 2)
    fail("Both groups need at least 2 non-missing observations.")

  n1 <- sum(!is.na(x1)); n2 <- sum(!is.na(x2))
  m1 <- mean(x1, na.rm = TRUE); m2 <- mean(x2, na.rm = TRUE)
  s1 <- stats::sd(x1, na.rm = TRUE); s2 <- stats::sd(x2, na.rm = TRUE)
```

```

dfree <- n1 + n2 - 2
sp2 <- (((n1 - 1) * (s1^2)) + ((n2 - 1) * (s2^2))) / dfree
sp  <- sqrt(sp2)

tstar <- qt(0.975, dfree)
se <- sp * sqrt(1/n1 + 1/n2)
exp_ci <- c((m1 - m2) - tstar*se, (m1 - m2) + tstar*se)

x <- .result
if (!is.numeric(x) || length(x) != 2L || any(!is.finite(x)))
fail("Print a length-2 numeric vector: c(lower, upper).")

if (max(abs(x - exp_ci)) < 1e-6)
pass(" Correct pooled 95% CI for FiveStar - NotFive (best_score).")
else
fail("Not quite - recheck pooled s_p, df, and t* with stars==5 vs stars!=5.")
})

```

## Q2 — Pooled (Equal Variances): Using `t.test()`

Compute a 90% two-sided confidence interval for

$$\mu_{Adelie} - \mu_{Gentoo}$$

for bill length (mm) using pooled variances. Use the local file penguins.csv and only the two species Adelie and Gentoo.

Photo by Adam Tarshis on Unsplash

### Info

For equal-variance two-sample CIs, `t.test(yield ~ method, var.equal = TRUE, conf.level = 0.90)` uses the pooled standard error and  $df = n_1 + n_2 - 2$ . To get the CI for Adelie – Gentoo, set the factor level order to `c("Adelie", "Gentoo")`.

### Preview

Run this code chunk to get a glimpse of the dataset. Feel free to change the values to visualize more/less number of rows.

```
#| echo: true
df <- read.csv("penguins.csv")
head(df[, c("species", "bill_length_mm", "island", "sex")], 8)
df2 <- subset(df, species %in% c("Adelie", "Gentoo") & is.finite(bill_length_mm))
set.seed(258)
grp_idx <- split(seq_len(nrow(df2)), df2$species)
idx <- unlist(lapply(grp_idx, function(i) sample(i, min(120, length(i)))))
df3 <- df2[idx, ]
par(mar = c(4, 4, 1, 1))
boxplot(bill_length_mm ~ species, data = df3, notch = FALSE, outline = FALSE,
ylab = "bill_length_mm", xlab = "Species")
```

### i Info

When  $\sigma$  is known,

$$\text{CI}_{1-\alpha} : \bar{x} \pm x_{1-\alpha/2} \frac{\sigma}{\sqrt{n}}.$$

```
#| exercise: q2_pooled_penguins
#| exercise.lines: 8
#| echo: false

df <- read.csv("penguins.csv")
keep <- df$species %in% c(" ", " ") #filter based on species needed
df <- subset(df, keep & !is.na( ))

df$method <- factor(df$ , levels = c(" ", " ")) # order matters for the contrast
df$yield <- df$bill_length_mm

tt <- t.test( ~ , data = df, var.equal = TRUE, conf.level = )
tt$conf.int #gives confidence interval
```

Use a two-group factor for the species and the pooled-variance option in `t.test`. Ensure you're returning just the interval bounds, not the whole test object.

*Solution.*

```
#| exercise: q2_pooled_penguins
#| solution: true
```

```

df <- read.csv("penguins.csv")
keep <- df$species %in% c("Adelie", "Gentoo")
df <- subset(df, keep & !is.na(bill_length_mm))

df$method <- factor(df$species, levels = c("Adelie", "Gentoo")) # order matters for the contrast
df$yield <- df$bill_length_mm

tt <- t.test(yield ~ method, data = df, var.equal = TRUE, conf.level = 0.90)
tt$conf.int

```

  

```

#| exercise: q2_pooled_penguins
#| check: true
gradethis::grade_this({
  df <- tryCatch(read.csv("penguins.csv"), error = function(e) NULL)
  if (is.null(df)) fail("Couldn't read 'penguins.csv'.")

  need <- c("species", "bill_length_mm")
  if (!all(need %in% names(df))) fail("CSV must have 'species' and 'bill_length_mm' columns.")

  keep <- df$species %in% c("Adelie", "Gentoo")
  df <- subset(df, keep & !is.na(bill_length_mm))
  if (nrow(df) < 4) fail("Need data for both Adelie and Gentoo with non-missing bill_length_mm")

  df$method <- factor(df$species, levels = c("Adelie", "Gentoo"))
  df$yield <- df$bill_length_mm

  exp_ci <- t.test(yield ~ method, data = df, var.equal = TRUE, conf.level = 0.90)$conf.int
})

```

```

x <- .result
if (!is.numeric(x) || length(x) != 2L || any(!is.finite(x)))
fail("Print a 2-number numeric vector: c(lower, upper).")

if (max(abs(x - exp_ci)) < 1e-6)
pass(" Correct 90% pooled CI for Adelie - Gentoo (bill length).")
else
fail("Not quite - check species filtering, factor level order, var.equal=TRUE, and conf.level")
})

```

### Q3 — Welch (Unequal Variances): Using `t.test()`

Compute an 80% two-sided confidence interval for

$$\mu_{Chinstrap} - \mu_{Gentoo}$$

for flipper\_length\_mm using Welch's unequal-variance method (the default in `t.test`). Use the local file `penguins.csv` and only the species Chinstrap and Gentoo. Print only the 2-number CI vector.

#### Info

`t.test(y ~ g, conf.level = 0.80)` uses Welch by default (`var.equal = FALSE`). To obtain Chinstrap – Gentoo, set the factor level order to `c("Chinstrap", "Gentoo")`.

#### Preview

Run this code chunk to get a glimpse of the dataset. Feel free to change the values to visualize more/less number of rows.

```
#| echo: true
df <- read.csv("penguins.csv")
head(df[, c("species", "bill_length_mm", "island", "sex")], 8)
df2 <- subset(df, species %in% c("Chinstrap", "Gentoo") & is.finite(flipper_length_mm))
set.seed(258)
grp_idx <- split(seq_len(nrow(df2)), df2$species)
idx <- unlist(lapply(grp_idx, function(i) sample(i, min(120, length(i)))))
df3 <- df2[idx, ]
par(mar = c(4, 4, 1, 1))
boxplot(flipper_length_mm ~ species, data = df3, notch = FALSE, outline = FALSE,
ylab = "flipper_length_mm", xlab = "Species")
```

```
#| exercise: q4_welch_ttest_penguins
#| exercise.lines: 8
#| echo: false
df <- read.csv("penguins.csv")

keep <- #filter based on species needed here
df <- subset(df, keep & !is.na( ))

df$method <- factor(df$ , levels = c(" ", " "))
df$yield <-

tt <- t.test( ~ , data = df, conf.level = ) # Welch by default
```

Ensure you keep exactly two species and set their order to match the contrast. Return just the two CI bounds from the test result.

*Solution.*

```
#| exercise: q4_welch_ttest_penguins
#| solution: true

df <- read.csv("penguins.csv")

keep <- df$species %in% c("Chinstrap", "Gentoo")
df <- subset(df, keep & !is.na(flipper_length_mm))

df$method <- factor(df$species, levels = c("Chinstrap", "Gentoo"))
df$y <- df$flipper_length_mm
```

```

tt <- t.test(y ~ method, data = df, conf.level = 0.80) # Welch by default
tt$conf.int

#| exercise: q4_welch_ttest_penguins
#| check: true
gradethis::grade_this({
df <- tryCatch(read.csv("penguins.csv"), error = function(e) NULL)
if (is.null(df)) fail("Couldn't read 'penguins.csv'.")
need <- c("species", "flipper_length_mm")
if (!all(need %in% names(df))) fail("CSV must have 'species' and 'flipper_length_mm' columns")

keep <- df$species %in% c("Chinstrap", "Gentoo")
df <- subset(df, keep & !is.na(flipper_length_mm))
if (nrow(df) < 4) fail("Need data for both Chinstrap and Gentoo with non-missing flipper_length_mm")

df$method <- factor(df$species, levels = c("Chinstrap", "Gentoo"))
df$y <- df$flipper_length_mm

exp_ci <- t.test(y ~ method, data = df, conf.level = 0.80)$conf.int

x <- .result
if (!is.numeric(x) || length(x) != 2L || any(!is.finite(x)))
fail("Print a 2-number numeric vector: c(lower, upper).")

if (max(abs(x - exp_ci)) < 1e-6)
pass(" Correct 80% Welch CI for Chinstrap - Gentoo (flipper length).")
else
fail("Not quite - check species filter, factor order, and conf.level = 0.80.")
})

```

## Q4 — Two-Sample CI for a Difference of Proportions

In recipe\_reviews.csv, compare the proportion of reviews where thumbs\_up > thumbs\_down between FiveStar (stars == 5) and NotFive (stars != 5). Compute a 95% two-sided CI for

$$p_{FiveStar} - p_{NotFive}$$

- . Print c(lower, upper).

## i Preview

Run this code chunk to get a glimpse of the dataset. Feel free to change the values to visualize more/less number of rows.

```
#| echo: true
df <- read.csv("recipe_reviews.csv")
head(df[, c("stars", "thumbs_up", "thumbs_down", "best_score")], 10)
df2 <- subset(df, is.finite(thumbs_up) & is.finite(thumbs_down) & is.finite(stars))
df2$group <- ifelse(df2$stars == 5, "FiveStar", "NotFive")
df2$success <- as.integer(df2$thumbs_up > df2$thumbs_down)
set.seed(258)
grp_idx <- split(seq_len(nrow(df2)), df2$group)
idx <- unlist(lapply(grp_idx, function(i) sample(i, min(400, length(i)))))
df3 <- df2[idx, ]
p <- tapply(df3$success, df3$group, function(x) mean(x, na.rm = TRUE))
par(mar = c(4,4,1,1))
barplot(p, ylim = c(0,1), ylab = "Proportion (thumbs_up > thumbs_down)", xlab = "Group")
```

## i Info

CI for difference of proportions:

$$\text{CI}_{1-\alpha} : \hat{p}_1 - \hat{p}_2 \pm z_{1-\alpha/2} * \sqrt{\frac{\hat{p}_1(1-\hat{p}_1)}{n_1} + \frac{\hat{p}_2(1-\hat{p}_2)}{n_2}}.$$

```
#| exercise: q11_wald_2prop_reviews
#| exercise.lines: 10
#| echo: false
df <- read.csv(" ")
g1 <- subset(df, stars == 5) #segregating the dataset
g2 <- subset(df, stars != 5)

count1 <- as.integer(g1$thumbs_up > g1$thumbs_down) #helps in making 0/1 success.
#Makes all thumbs_up values 1 and thumbs_down values 0.
count2 <- as.integer(g2$thumbs_up > g2$thumbs_down)

n1 <- sum(is.finite( )); n2 <- sum(is.finite( )) #is.finite helps to
#filter only usable rows
p1 <- mean( , na.rm = TRUE)
p2 <- mean( , na.rm = TRUE)
```

```
se <-
z <-
```

Compute

$$\hat{p}$$

per group, then compute their SE and z\* at 0.975.

*Solution.*

```
#| exercise: q11_wald_2prop_reviews
#| solution: true
df <- read.csv("recipe_reviews.csv")
g1 <- subset(df, stars == 5) #segregating the dataset
g2 <- subset(df, stars != 5)
count1 <- as.integer(g1$thumbs_up > g1$thumbs_down)
count2 <- as.integer(g2$thumbs_up > g2$thumbs_down)
n1 <- sum(is.finite(count1)); n2 <- sum(is.finite(count2))
p1 <- mean(count1, na.rm = TRUE); p2 <- mean(count2, na.rm = TRUE)
se <- sqrt(p1*(1-p1)/n1 + p2*(1-p2)/n2)
z <- qnorm(0.975);
(p1 - p2) + c(-1,1)*z*se
```

```
#| exercise: q11_wald_2prop_reviews
#| check: true
gradethis::grade_this({
df <- tryCatch(read.csv("recipe_reviews.csv"), error = function(e) NULL)
if (is.null(df)) fail("Couldn't read 'recipe_reviews.csv'.")
g1 <- subset(df, stars == 5)
g2 <- subset(df, stars != 5)
y1 <- as.integer(g1$thumbs_up > g1$thumbs_down)
y2 <- as.integer(g2$thumbs_up > g2$thumbs_down)
n1 <- sum(is.finite(y1)); n2 <- sum(is.finite(y2))
if (min(n1, n2) < 2) fail("Both groups need at least 2 finite rows.")
p1 <- mean(y1, na.rm = TRUE); p2 <- mean(y2, na.rm = TRUE)
se <- sqrt(p1*(1-p1)/n1 + p2*(1-p2)/n2)
z <- qnorm(0.975); est <- p1 - p2
exp <- est + c(-1,1)*z*se
x <- .result
ok <- is.numeric(x) && length(x)==2L && all(is.finite(x))
if (!ok) fail("Print c(lower, upper).")
if (max(abs(x - exp)) < 1e-6) pass(" Correct 95% Wald CI for p_FiveStar - p_NotFive.")
```

```
else fail("Recheck successes, n1/n2, and z*.")
})
```

## Q5 — Two-Sample CI for a Difference of Proportions using prop.test

In penguins.csv, compare the male proportion between Chinstrap and Gentoo. Compute an evidence-only CI (no hypothesis interpretation needed) for  $p_{Chinstrap} - p_{Gentoo}$  at 90% and print c(lower upper).

### i Preview

```
#| echo: true
df <- read.csv("penguins.csv")
head(df[, c("species", "sex", "island")], 10)
df2 <- subset(df, species %in% c("Chinstrap", "Gentoo") & !is.na(sex))
set.seed(258)
grp_idx <- split(seq_len(nrow(df2)), df2$species)
idx <- unlist(lapply(grp_idx, function(i) sample(i, min(200, length(i)))))
df3 <- df2[idx, ]
p <- with(df3, tapply(sex == "male", species, function(x) mean(x, na.rm = TRUE)))
par(mar = c(4,4,1,1))
barplot(p, ylim = c(0,1), ylab = "Proportion male", xlab = "Species")
```

### i Info

The `prop.test()` function in R performs tests of proportions, allowing for the comparison of proportions across groups or against specific values. Here is an example according to this question. here's a super simple one-sample `prop.test` using `penguins.csv`. It asks: among Gentoo penguins, what fraction are male? It prints the estimate and a 95% CI. Run the code to see!

```
#| echo: true

# One-sample prop.test example: proportion of males within Gentoo

df <- read.csv("penguins.csv")

g <- subset(df, species == "Gentoo" & !is.na(sex))
x <- sum(g$sex == "male")    # number of "successes"
n <- nrow(g)                 # total trials

prop.test(x = x, n = n, conf.level = 0.95, correct = FALSE)
```

```
#| exercise: q12_prop_test_penguins
#| exercise.lines: 10
#| echo: false
df <- read.csv("penguins.csv")
df <- subset(df, species %in% c(" ", " ") & !is.na(sex)) #filtering based
#on the species we need and removing NA rows
x1 <- sum(df$species=="Chinstrap" & df$sex==" ")
n1 <-
x2 <- sum(df$species=="Gentoo" & df$sex==" ")
n2 <-
prop.test(x = c(   ,   ), n = c(   ,   ), conf.level =   , correct = FALSE)$conf.int
```

Filter to two species, count “male” in each, and feed (x1,n1),(x2,n2) to prop.test.

*Solution.*

```
#| exercise: q12_prop_test_penguins
#| solution: true
df <- read.csv("penguins.csv")
df <- subset(df, species %in% c("Chinstrap","Gentoo") & !is.na(sex))
x1 <- sum(df$species=="Chinstrap" & df$sex=="male")
n1 <- sum(df$species=="Chinstrap")
x2 <- sum(df$species=="Gentoo" & df$sex=="male")
n2 <- sum(df$species=="Gentoo")
prop.test(x = c(x1, x2), n = c(n1, n2), conf.level = 0.90, correct = FALSE)$conf.int
```

```
#| exercise: q12_prop_test_penguins
#| check: true
gradethis::grade_this({
```

```

df <- tryCatch(read.csv("penguins.csv"), error=function(e) NULL)
if (is.null(df)) fail("Couldn't read 'penguins.csv'.")
df <- subset(df, species %in% c("Chinstrap", "Gentoo") & !is.na(sex))
x1 <- sum(df$species=="Chinstrap" & df$sex=="male"); n1 <- sum(df$species=="Chinstrap")
x2 <- sum(df$species=="Gentoo" & df$sex=="male"); n2 <- sum(df$species=="Gentoo")
exp <- prop.test(x=c(x1,x2), n=c(n1,n2), conf.level=0.90, correct=FALSE)$conf.int
x <- .result
ok <- is.numeric(x) && length(x)==2L && all(is.finite(x))
if (!ok) fail("Print c(lower, upper).")
if (max(abs(x - exp)) < 1e-6) pass(" Correct 90% CI for p_Chinstrap - p_Gentoo.")
else fail("Check species filter, male counts, and conf.level=0.90.")
})

```

## Q6 — Two-Sample CI for a Difference of Proportions (island=Biscoe)

Using penguins.csv, compare the proportion of penguins on Biscoe island between Adelie and Gentoo. Compute a 95% two-sided CI for  $p_{Adelie} - p_{Gentoo}$  and print c(lower upper).

Preview

```

#| echo: true
df <- read.csv("penguins.csv")
head(df[, c("species", "sex", "island")], 10)
df2 <- subset(df, species %in% c("Adelie", "Gentoo") & !is.na(island))
set.seed(258)
grp_idx <- split(seq_len(nrow(df2)), df2$species)
idx <- unlist(lapply(grp_idx, function(i) sample(i, min(200, length(i)))))
df3 <- df2[idx, ]
p <- with(df3, tapply(island == "Biscoe", species, function(x) mean(x, na.rm = TRUE)))
par(mar = c(4,4,1,1))
barplot(p, ylim = c(0,1), ylab = 'Proportion on "Biscoe"', xlab = "Species")

```

```

#| exercise: q5_diff_prop_biscoe
#| exercise.lines: 6
#| echo: false
df <- read.csv("penguins.csv")
sub <- subset(df, species %in% c(" ", " ") & !is.na( )) #filter based on species

x1 <- sum(sub$species==" " & sub$ == " ") #add additional filter for island
n1 <- sum(sub$species=="Adelie")
x2 <- sum(sub$species==" " & sub$ == " ")

```

```

n2 <- sum(sub$species=="Gentoo")

prop.test(x=c( , ), n=c( , ), conf.level= , correct=FALSE)$conf.int

```

Count “Biscoe” within each species to get x1, x2. Use totals for n1, n2, then prop.test(...)\$conf.int. Order your counts as (Adelie, Gentoo) to get Adelie – Gentoo.

*Solution.*

```

#| exercise: q5_diff_prop_biscoe
#| solution: true
df <- read.csv("penguins.csv")
sub <- subset(df, species %in% c("Adelie", "Gentoo") & !is.na(island))
x1 <- sum(sub$species=="Adelie" & sub$island=="Biscoe"); n1 <- sum(sub$species=="Adelie")
x2 <- sum(sub$species=="Gentoo" & sub$island=="Biscoe"); n2 <- sum(sub$species=="Gentoo")
prop.test(x=c(x1,x2), n=c(n1,n2), conf.level=0.95, correct=FALSE)$conf.int

```

```

#| exercise: q5_diff_prop_biscoe
#| check: true
gradethis::grade_this({
  df <- tryCatch(read.csv("penguins.csv"), error=function(e) NULL)
  if (is.null(df)) fail("Couldn't read 'penguins.csv'.")
  sub <- subset(df, species %in% c("Adelie", "Gentoo") & !is.na(island))
  x1 <- sum(sub$species=="Adelie" & sub$island=="Biscoe"); n1 <- sum(sub$species=="Adelie")
  x2 <- sum(sub$species=="Gentoo" & sub$island=="Biscoe"); n2 <- sum(sub$species=="Gentoo")
  exp <- prop.test(x=c(x1,x2), n=c(n1,n2), conf.level=0.95, correct=FALSE)$conf.int
  r <- .result
  ok <- is.numeric(r) && length(r)==2L && all(is.finite(r))
  if (!ok) fail("Print a numeric vector c(lower, upper).")
  if (max(abs(r - exp)) < 1e-6) pass(" Correct 95% CI for p_Adelie - p_Gentoo (Biscoe).")
  else fail("Recheck counts, totals, order (Adelie first), and conf.level=0.95.")
})

```

## Q7 — Two-Sample CI for a Ratio of Variances (manual F-based)

Using recipe\_reviews.csv, compare the variance of best\_score for FiveStar vs NotFive. Build a 95% CI for  $\sigma_{FiveStar}^2/\sigma_{NotFive}^2$  and print c(lower, upper).

## i Preview

Here is a glimpse of the dataset:

```
#| echo: true
df <- read.csv("recipe_reviews.csv")
head(df[, c("stars", "best_score")], 10)
df2 <- subset(df, is.finite(best_score) & is.finite(stars))
df2$group <- ifelse(df2$stars == 5, "FiveStar", "NotFive")
set.seed(258)
grp_idx <- split(seq_len(nrow(df2)), df2$group)
idx <- unlist(lapply(grp_idx, function(i) sample(i, min(200, length(i)))))
df3 <- df2[idx, ]
ylim <- quantile(df3$best_score, c(.01, .99), na.rm = TRUE)
par(mar = c(4,4,1,1))
boxplot(best_score ~ group, data = df3, notch = FALSE, outline = FALSE, ylim = ylim,
ylab = "best_score", xlab = "Group")
```

## i Info

Taking group 1 as FiveStar here, with  $s_1^2$  and  $s_2^2$  as sample variances, here is the CI:

$$\text{CI}_{1-\alpha}\left(\frac{\sigma_1^2}{\sigma_2^2}\right) = \left(\frac{s_1^2/s_2^2}{F_{1-\alpha/2, df_1, df_2}}, \frac{s_1^2/s_2^2}{F_{\alpha/2, df_1, df_2}}\right), \quad df_1 = n_1 - 1, df_2 = n_2 - 1.$$

```
#| exercise: q7_ratio_var_reviews
#| exercise.lines: 12
#| echo: false
df <- read.csv("recipe_reviews.csv")
x1 <- subset(df, stars == 5)$
x2 <- subset(df, stars != 5)$
x1 <- x1[is.finite(x1)] #only keep usable rows
x2 <- x2[is.finite(x2)]
n1 <- length( ); n2 <- length( )
s1 <- stats:: (x1); s2 <- stats:: (x2) #computing variance of both vectors x1,x2
df1 <- ; df2 <-
ratio <- s1/s2
F_lo <-
F_hi <-
```

Use FiveStar as numerator. Compute  $s_1^2$  and  $s_2^2$  and divide by the F cutoffs (0.975 and 0.025).

*Solution.*

```
#| exercise: q7_ratio_var_reviews
#| solution: true
df <- read.csv("recipe_reviews.csv")
x1 <- subset(df, stars == 5)$best_score
x2 <- subset(df, stars != 5)$best_score
x1 <- x1[is.finite(x1)]
x2 <- x2[is.finite(x2)]
n1 <- length(x1)
n2 <- length(x2)
s1 <- stats::var(x1)
s2 <- stats::var(x2)
df1 <- n1 - 1
df2 <- n2 - 1
ratio <- s1/s2
F_lo <- qf(0.975, df1, df2)
F_hi <- qf(0.025, df1, df2)
c(ratio/F_lo, ratio/F_hi)
```

```
#| exercise: q7_ratio_var_reviews
#| check: true
gradethis::grade_this({
  df <- tryCatch(read.csv("recipe_reviews.csv"), error=function(e) NULL)
  if (is.null(df)) fail("Couldn't read 'recipe_reviews.csv'.")
  x1 <- subset(df, stars == 5)$best_score
  x2 <- subset(df, stars != 5)$best_score
  x1 <- x1[is.finite(x1)]; x2 <- x2[is.finite(x2)]
  if (min(length(x1), length(x2)) < 3) fail("Each group needs at least 3 finite rows.")
  s1 <- stats::var(x1); s2 <- stats::var(x2)
  df1 <- length(x1)-1; df2 <- length(x2)-1
  ratio <- s1/s2
  F_lo <- qf(0.975, df1, df2); F_hi <- qf(0.025, df1, df2)
  exp <- c(ratio/F_lo, ratio/F_hi)
  x <- .result
  ok <- is.numeric(x) && length(x)==2L && all(is.finite(x))
  if (!ok) fail("Print c(lower, upper).")
  if (max(abs(x - exp)) < 1e-6) pass(" Correct 95% CI for `FiveStar` / `NotFive.`")
  else fail("Check df and F quantiles (0.975, 0.025).")
})
```

## Q8 - Two-Sample CI for a Ratio of Variances using var.test

In penguins.csv, compare variance of bill\_length\_mm between Adelie (group 1) and Gentoo (group 2). Compute an 80% CI for  $\sigma_{Adelie}^2/\sigma_{Gentoo}^2$  using var.test and print the two bounds.

### i Info

Here is an example to demonstrate how var.test works. We are taking the recipe reviews dataset to showcase the usage. Run the code to view results.

```
#thumbs_up variance: FiveStar vs NotFive (95% CI)
df <- read.csv("recipe_reviews.csv")

x1 <- df$thumbs_up[df$stars == 5]
x2 <- df$thumbs_up[df$stars != 5]
x1 <- x1[is.finite(x1)]; x2 <- x2[is.finite(x2)]

vt <- var.test(x1, x2, conf.level = 0.95)

vt$estimate # s_FiveStar^2 / s_NotFive^2
vt$conf.int # 95% CI for _FiveStar^2 / _NotFive^2
```

### i Preview

```
#| echo: true
df <- read.csv("penguins.csv")
head(df[, c("species", "bill_length_mm")], 10)
df2 <- subset(df, species %in% c("Adelie", "Gentoo") & is.finite(bill_length_mm))
set.seed(258)
grp_idx <- split(seq_len(nrow(df2)), df2$species)
idx <- unlist(lapply(grp_idx, function(i) sample(i, min(120, length(i)))))
df3 <- df2[idx, ]
par(mar = c(4, 4, 1, 1))
boxplot(bill_length_mm ~ species, data = df3, notch = FALSE, outline = FALSE,
ylab = "bill_length_mm", xlab = "Species")
```

```
#| exercise: q8_var_test_penguins
#| exercise.lines: 8
#| echo: false
df <- read.csv("penguins.csv")
df <- subset(df, species %in% c( " ", " " ) & !is.na( ))
x <- df$ [df$ == "Adelie"]
```

```
y <- df$ [df$ == "Gentoo"]
var.test( )$conf.int
```

Filter to Adelie and Gentoo, remove NAs, then feed the two numeric vectors to var.test.

*Solution.*

```
#| exercise: q8_var_test_penguins
#| solution: true
df <- read.csv("penguins.csv")
df <- subset(df, species %in% c("Adelie", "Gentoo") & !is.na(bill_length_mm))
x <- df$bill_length_mm[df$species=="Adelie"]
y <- df$bill_length_mm[df$species=="Gentoo"]
var.test(x, y, conf.level = 0.80)$conf.int
```

```
#| exercise: q8_var_test_penguins
#| check: true
gradethis::grade_this({
df <- tryCatch(read.csv("penguins.csv"), error=function(e) NULL)
if (is.null(df)) fail("Couldn't read 'penguins.csv'.")
df <- subset(df, species %in% c("Adelie", "Gentoo") & !is.na(bill_length_mm))
x <- df$bill_length_mm[df$species=="Adelie"]
y <- df$bill_length_mm[df$species=="Gentoo"]
exp <- var.test(x, y, conf.level = 0.80)$conf.int
z <- .result
ok <- is.numeric(z) && length(z)==2L && all(is.finite(z))
if (!ok) fail("Print c(lower, upper).")
if (max(abs(z - exp)) < 1e-6) pass(" Correct 80% CI for ^_Adelie / ^_Gentoo.")
else fail("Confirm vector order (Adelie first) and conf.level=0.80.")
})
```